

1. Write a program to implement the naïve Bayesian classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering few test data sets.

Aim :

To write a Python program to implement a Naïve Bayesian classifier for a sample training dataset stored in a .CSV file. The program should accurately classify test data sets and compute the classifier's accuracy

1. Import necessary libraries

```
import pandas as pd
from sklearn import tree
from sklearn.preprocessing import LabelEncoder
from sklearn.naive_bayes import GaussianNB
```

2. Load data from CSV

```
data = pd.read_csv('tennisdata.csv')
print("The first 5 values of data is :\n",data.head())
```

Output:

The first 5 values of data is :

	Outlook	Temperature	Humidity	Windy	PlayTennis
0	Sunny	Hot	High	False	No
1	Sunny	Hot	High	True	No
2	Overcast	Hot	High	False	Yes
3	Rainy	Mild	High	False	Yes
4	Rainy	Cool	Normal	False	Yes

3. Obtain Train data and Train output

```
X = data.iloc[:, :-1]
print("\nThe First 5 values of train data is\n",X.head())
```

Output:

The First 5 values of train data is

	Outlook	Temperature	Humidity	Windy
0	Sunny	Hot	High	False
1	Sunny	Hot	High	True
2	Overcast	Hot	High	False
3	Rainy	Mild	High	False
4	Rainy	Cool	Normal	False

```
y = data.iloc[:, -1]
print("\nThe first 5 values of Train output is\n",y.head())
```

Output:

The first 5 values of Train output is

0	No
---	----

```
1 No
2 Yes
3 Yes
4 Yes
```

Name: PlayTennis, dtype: object

4. Convert then in numbers

```
le_outlook = LabelEncoder()
X.Outlook = le_outlook.fit_transform(X.Outlook)
```

```
le_Temperature = LabelEncoder()
X.Temperature = le_Temperature.fit_transform(X.Temperature)
```

```
le_Humidity = LabelEncoder()
X.Humidity = le_Humidity.fit_transform(X.Humidity)
```

```
le_Windy = LabelEncoder()
X.Windy = le_Windy.fit_transform(X.Windy)
```

```
print("\nNow the Train data is :\n",X.head())
```

Output:

Now the Train data is :

	Outlook	Temperature	Humidity	Windy
0	2	1	0	0
1	2	1	0	1
2	0	1	0	0
3	1	2	0	0
4	1	0	1	0

```
le_PlayTennis = LabelEncoder()
y = le_PlayTennis.fit_transform(y)
print("\nNow the Train output is\n",y)
```

Now the Train output is

```
[0 0 1 1 1 0 1 0 1 1 1 1 0]
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.20)
```

```
classifier = GaussianNB()
classifier.fit(X_train,y_train)
```

```
from sklearn.metrics import accuracy_score
print("Accuracy is:",accuracy_score(classifier.predict(X_test),y_test))
```

Output:

Accuracy is: 0.6666666666666666

CSV FILE

Outlook, Temperature, Humidity, Windy, PlayTennis

Sunny, Hot, High, False, No

Sunny, Hot, High, True, No

Overcast, Hot, High, False, Yes

Rainy, Mild, High, False, Yes

Rainy, Cool, Normal, False, Yes

Rainy, Cool, Normal, True, No

Overcast, Cool, Normal, True, Yes

Sunny, Mild, High, False, No

Sunny, Cool, Normal, False, Yes

Rainy, Mild, Normal, False, Yes

Sunny, Mild, Normal, True, Yes

Overcast, Mild, High, True, Yes

Overcast, Hot, Normal, False, Yes

Rainy, Mild, High, True, No

2. Implement Linear Regression in a given business scenario and comment on its efficiency and performance.

Aim :

To Implement Linear Regression in a specified business context to analyse its applicability and performance, evaluating its efficiency and predictive capability.

1. Import necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

2. Load the advertising dataset

```
data = pd.read_csv('Ex2.csv')
```

3. Display the first few rows of the dataset

```
print(data.head())
```

4. Prepare the data

```
X = data[['TV', 'Radio', 'Newspaper']]
y = data['Sales']
```

5. Split the data into training and test sets (80% training, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

6. Create a linear regression model

```
model = LinearRegression()
```

7. Train the model using the training sets

```
model.fit(X_train, y_train)
```

8. Make predictions using the testing set

```
y_pred = model.predict(X_test)
```

9. Evaluate the model

```
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
print(f'R^2 Score: {r2}')
```

10. Residual plot

```
residuals = y_test - y_pred
plt.figure(figsize=(8, 6))
plt.scatter(y_pred, residuals, color='blue')
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')
plt.show()
```

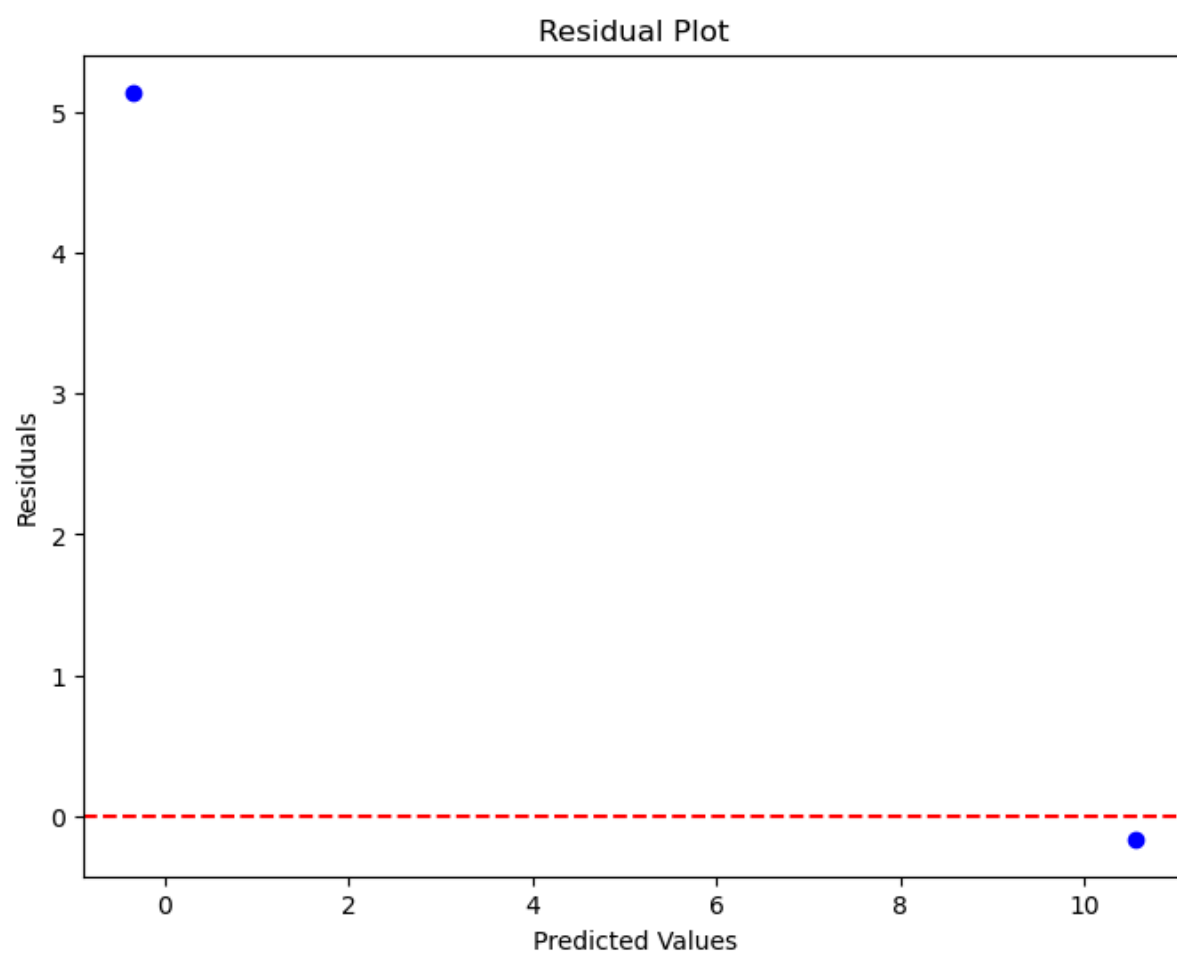
11. Predicted vs Actual plot


```
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, color='green')
plt.title('Predicted vs Actual')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--')
plt.show()
```

Output:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

Mean Squared Error: 13.201451790963432
R^2 Score: -0.6838586468065599





In []:

CSV File

TV, Radio, Newspaper, Sales

230.1, 37.8, 69.2, 22.1

44.5, 39.3, 45.1, 10.4

17.2, 45.9, 69.3, 9.3

151.5, 41.3, 58.5, 18.5

180.8, 10.8, 58.4, 12.9

8.7, 48.9, 75, 7.2

57.5, 32.8, 23.5, 11.8

120.2, 19.6, 11.6, 13.2

8.6, 2.1, 1, 4.8

199.8, 2.6, 21.2, 10.6

3. Implement SVM algorithm in a given business scenario and comment on its efficiency and performance.

Aim:

To Implement SVM algorithm in a given business scenario and comment on its efficiency and performance.

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
from sklearn.svm import SVC
```

1. Load dataset

```
data = pd.read_csv('Ex3.csv')
```

2. Handling missing values (if any)

```
data.fillna(method='ffill', inplace=True)
```

3. Encoding categorical variables

```
le = LabelEncoder()
data['Gender'] = le.fit_transform(data['Gender'])
```

4. Splitting the dataset into the Training set and Test set

```
X = data.drop('Purchase', axis=1)
y = data['Purchase']
```

5. Feature scaling

```
sc = StandardScaler()
X_scaled = sc.fit_transform(X)
```

6. Define SVM classifiers

```
classifiers = {
    'Linear SVM': SVC(kernel='linear', random_state=0),
    'Polynomial SVM': SVC(kernel='poly', degree=3, random_state=0),
    'RBF SVM': SVC(kernel='rbf', random_state=0)
}
```

7. Evaluate each classifier using k-fold cross-validation

```
for clf_name, clf in classifiers.items():
    scores = cross_val_score(clf, X_scaled, y, cv=5, scoring='accuracy')
    print(f"{clf_name} Cross-Validation Accuracy: {scores.mean():.2f} (+/- {scores.std() * 2:.2f})")
```


8. Train a selected SVM model (e.g., Linear SVM) on the entire dataset

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.25, random_state=0)
classifier = SVC(kernel='linear', random_state=0)
classifier.fit(X_train, y_train)
```

9. Make Predictions

```
y_pred = classifier.predict(X_test)
```

10. Performance Evaluation

```
cm = confusion_matrix(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
print(f"Performance Score: {f1:.2f}")
```

Output:

```
Linear SVM Cross-Validation Accuracy: 0.80 (+/- 0.33)
Polynomial SVM Cross-Validation Accuracy: 0.80 (+/- 0.53)
RBF SVM Cross-Validation Accuracy: 0.80 (+/- 0.53)
Performance Score: 1.00
```

CSV File

```
Age,Income,Gender,Purchase
25,50000,Male,0
45,64000,Female,1
35,58000,Female,0
50,72000,Male,1
23,48000,Male,0
31,52000,Female,0
46,60000,Female,1
29,55000,Male,0
52,75000,Male,1
48,68000,Female,1
36,59000,Male,0
28,53000,Female,0
27,52000,Female,0
44,61000,Male,1
33,57000,Female,1
```

4.Implement Decision Tree algorithm in a given business scenario and comment on its efficiency and performance.

Aim:

To Implement Decision Tree algorithm in a given business scenario and comment on its efficiency and performance.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score, f1_score
```

1.Load dataset

```
data = pd.read_csv('Ex4.csv')
```

2.Data preprocessing

```
# One-Hot Encoding for categorical variable 'Gender'
data = pd.get_dummies(data, columns=['Gender'], drop_first=True)
```

3.Verify columns after one-hot encoding

```
print(data.columns)
```

4.Splitting the dataset into the Training set and Test set

```
X = data.drop('TargetVariable', axis=1) # Adjust to your actual target
variable name
y = data['TargetVariable']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=0)
```

5.Building the Decision Tree model

```
classifier = DecisionTreeClassifier(random_state=0)
classifier.fit(X_train, y_train)
```

6.Making predictions

```
y_pred = classifier.predict(X_test)
```

7.Performance evaluation

```
cm = confusion_matrix(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

8.Outputting results

```
print("Confusion Matrix:\n", cm)
```

```
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1 Score: {f1:.2f}")
```

```
Index(['Age', 'Income', 'TargetVariable', 'Gender_Male'], dtype='object')
Confusion Matrix:
[[0 0]
 [1 3]]
Accuracy: 0.75
Precision: 1.00
Recall: 0.75
F1 Score: 0.86
```

CSV File

```
Age,Income,Gender,TargetVariable
25,50000,Male,0
45,64000,Female,1
35,58000,Female,0
50,72000,Male,1
23,48000,Male,0
31,52000,Female,0
46,60000,Female,1
29,55000,Male,0
52,75000,Male,1
48,68000,Female,1
36,59000,Male,0
28,53000,Female,0
27,52000,Female,0
44,61000,Male,1
33,57000,Female,1
```